ResNet-18 / CIFAR-10: LoRA recovery controls, adjacent-level faults, 16 clusters per tensor

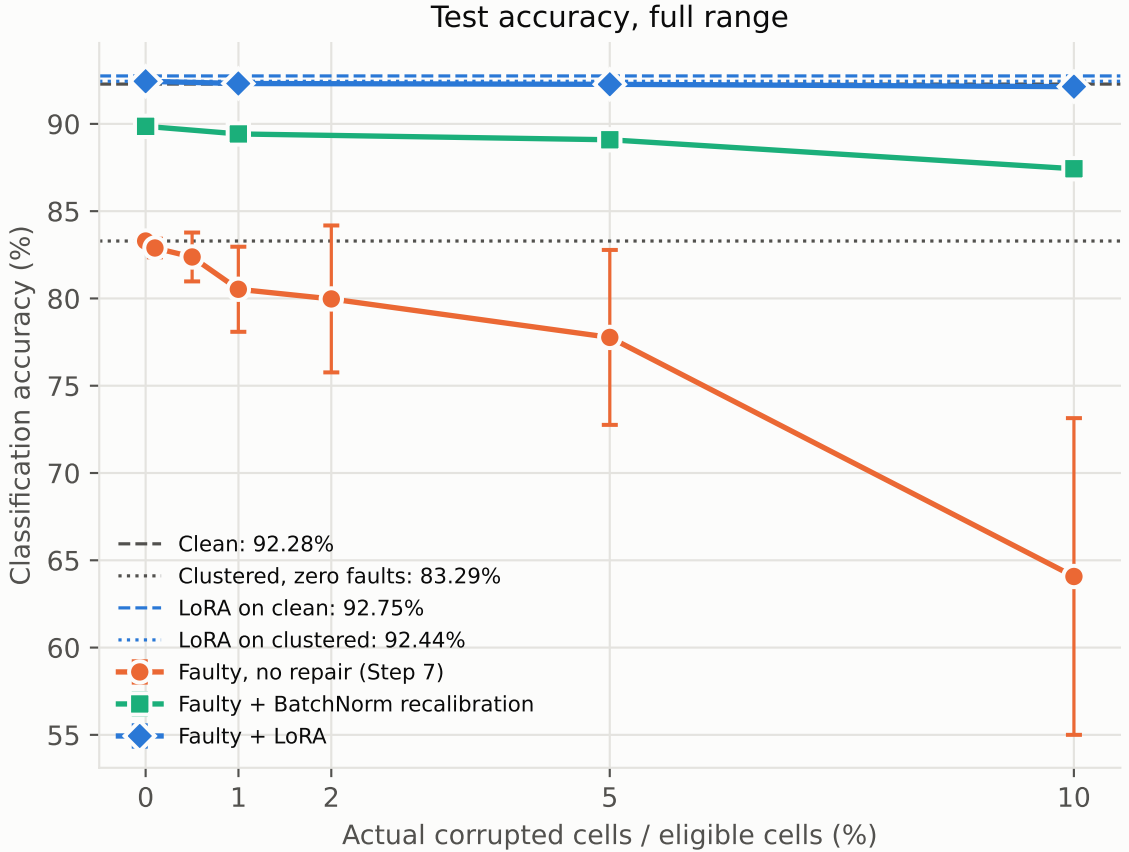

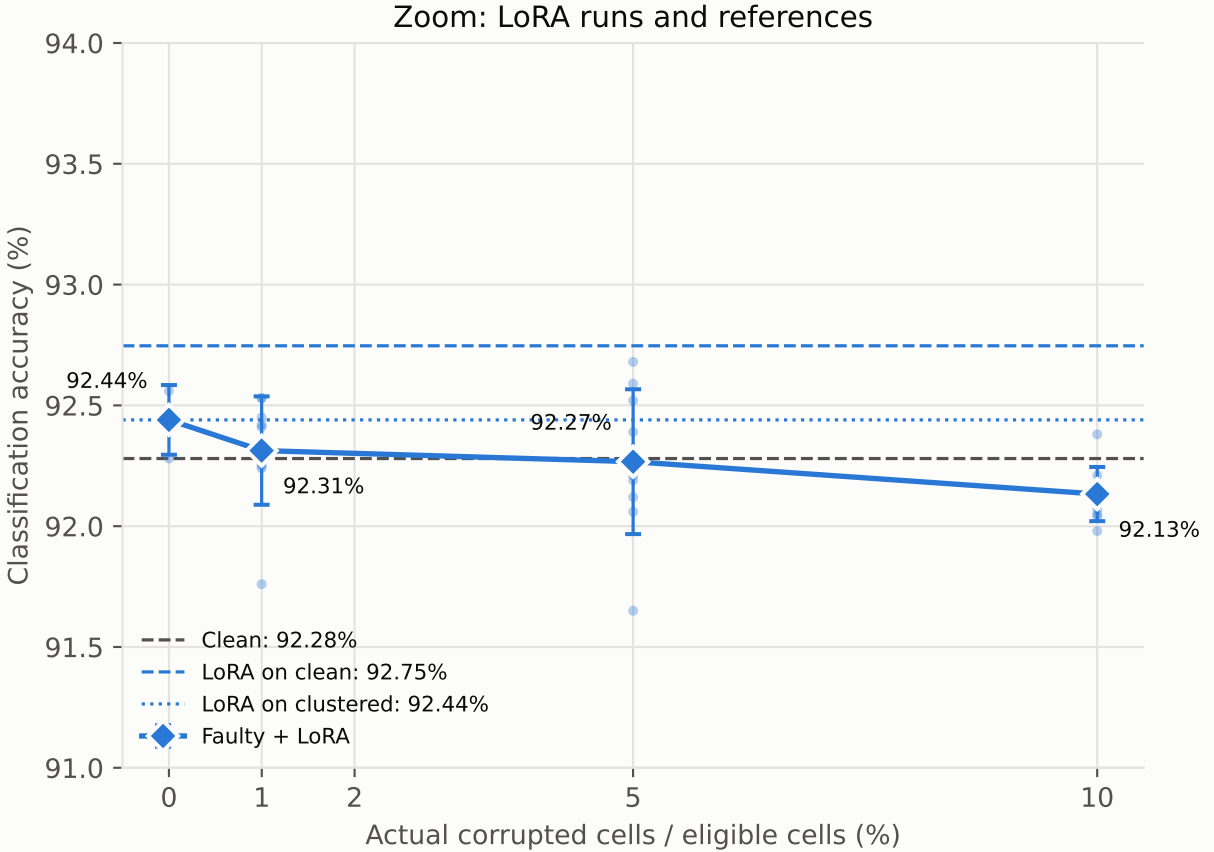

Means ± 1 sample SD. LoRA at 0%: clustered base, 3 training seeds; nonzero rates: 10 runs each (patterns × training seeds). Step 7 and BatchNorm recalibration: across patterns. Small dots: individual LoRA runs.