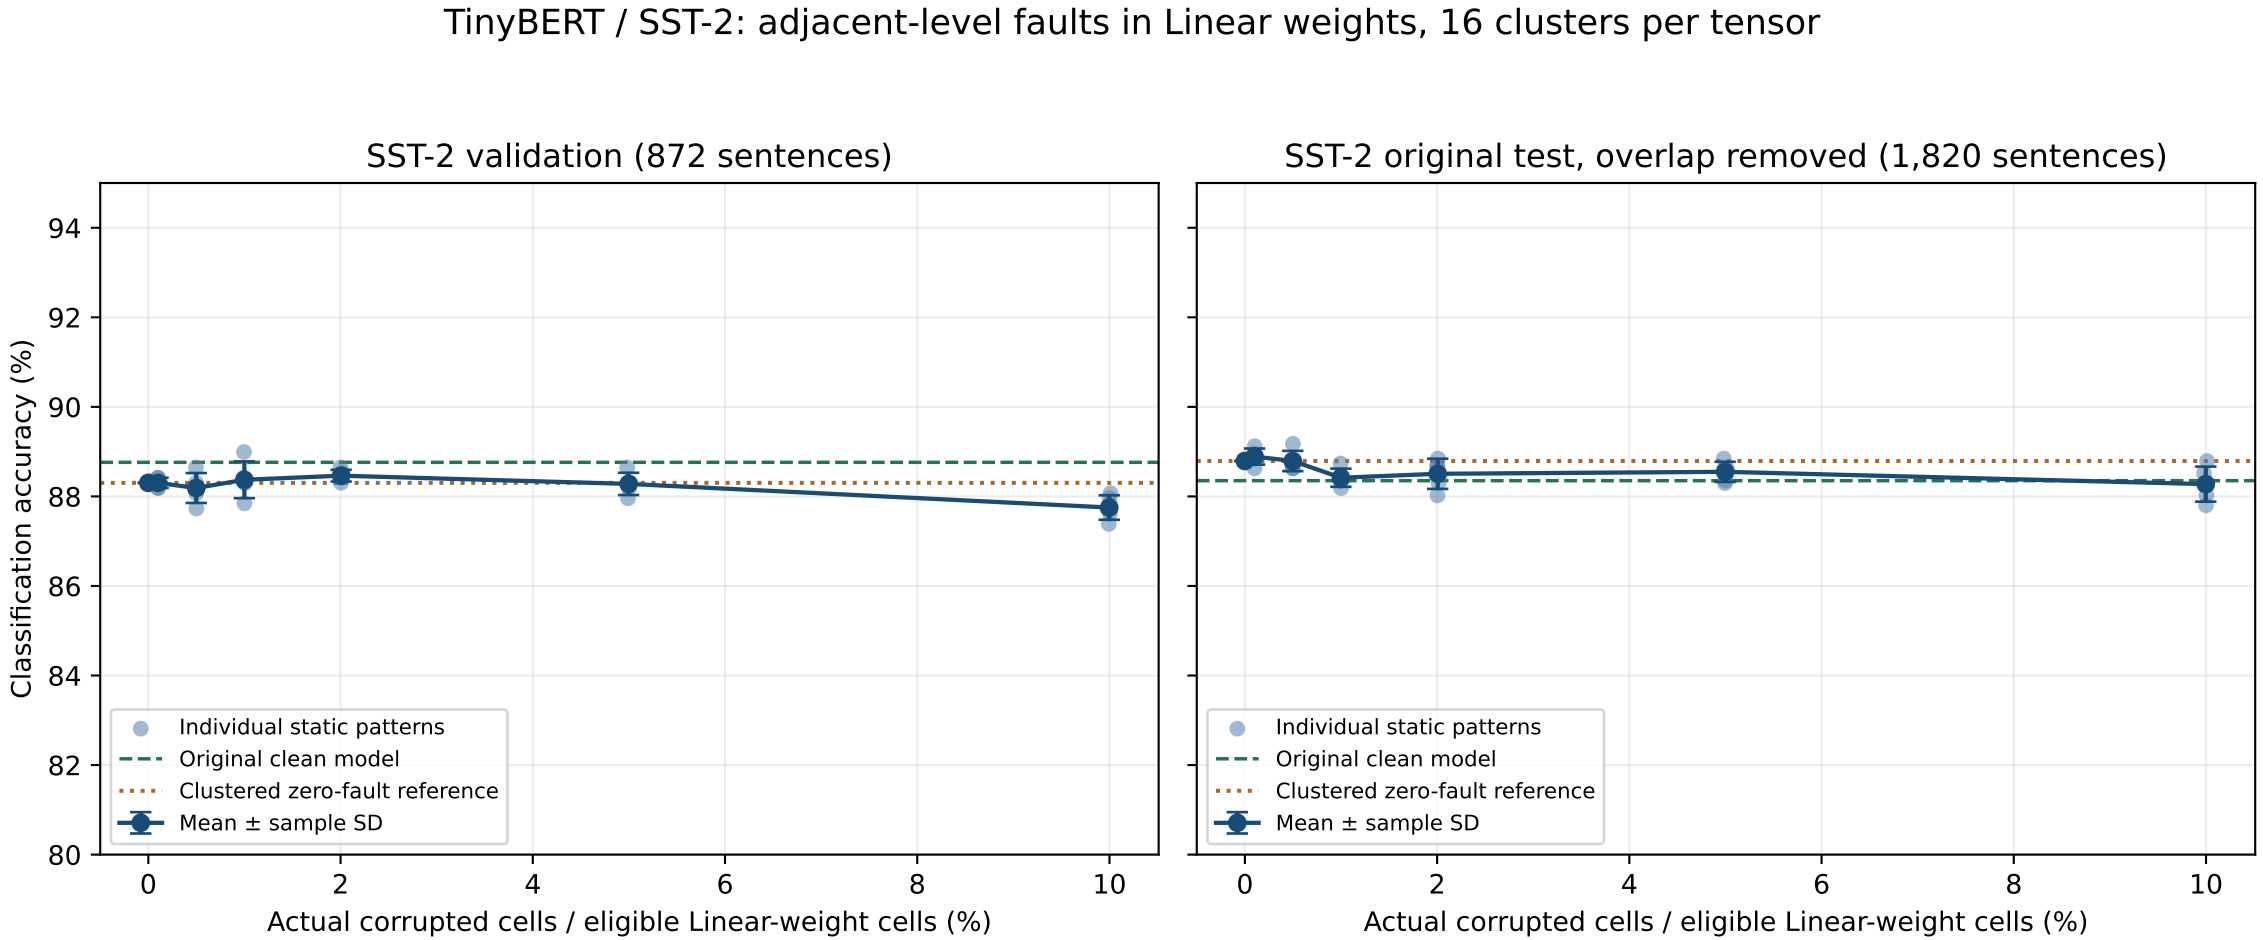
TinyBERT / SST-2: adjacent-level faults in Linear weights, 16 clusters per tensor
SST-2 validation (872 sentences)
SST-2 original test, overlap removed (1,820 sentences)
Classification accuracy (%)
94
92
90
88
86
84
82
80
0
2
4
6
8
10
Actual corrupted cells / eligible Linear-weight cells (%)
Individual static patterns
Original clean model
Clustered zero-fault reference
Mean ± sample SD
Nonzero rates: [5] patterns each; zero faults: one reference. Uniform selection; embeddings, LayerNorm, biases, codebooks reliable.