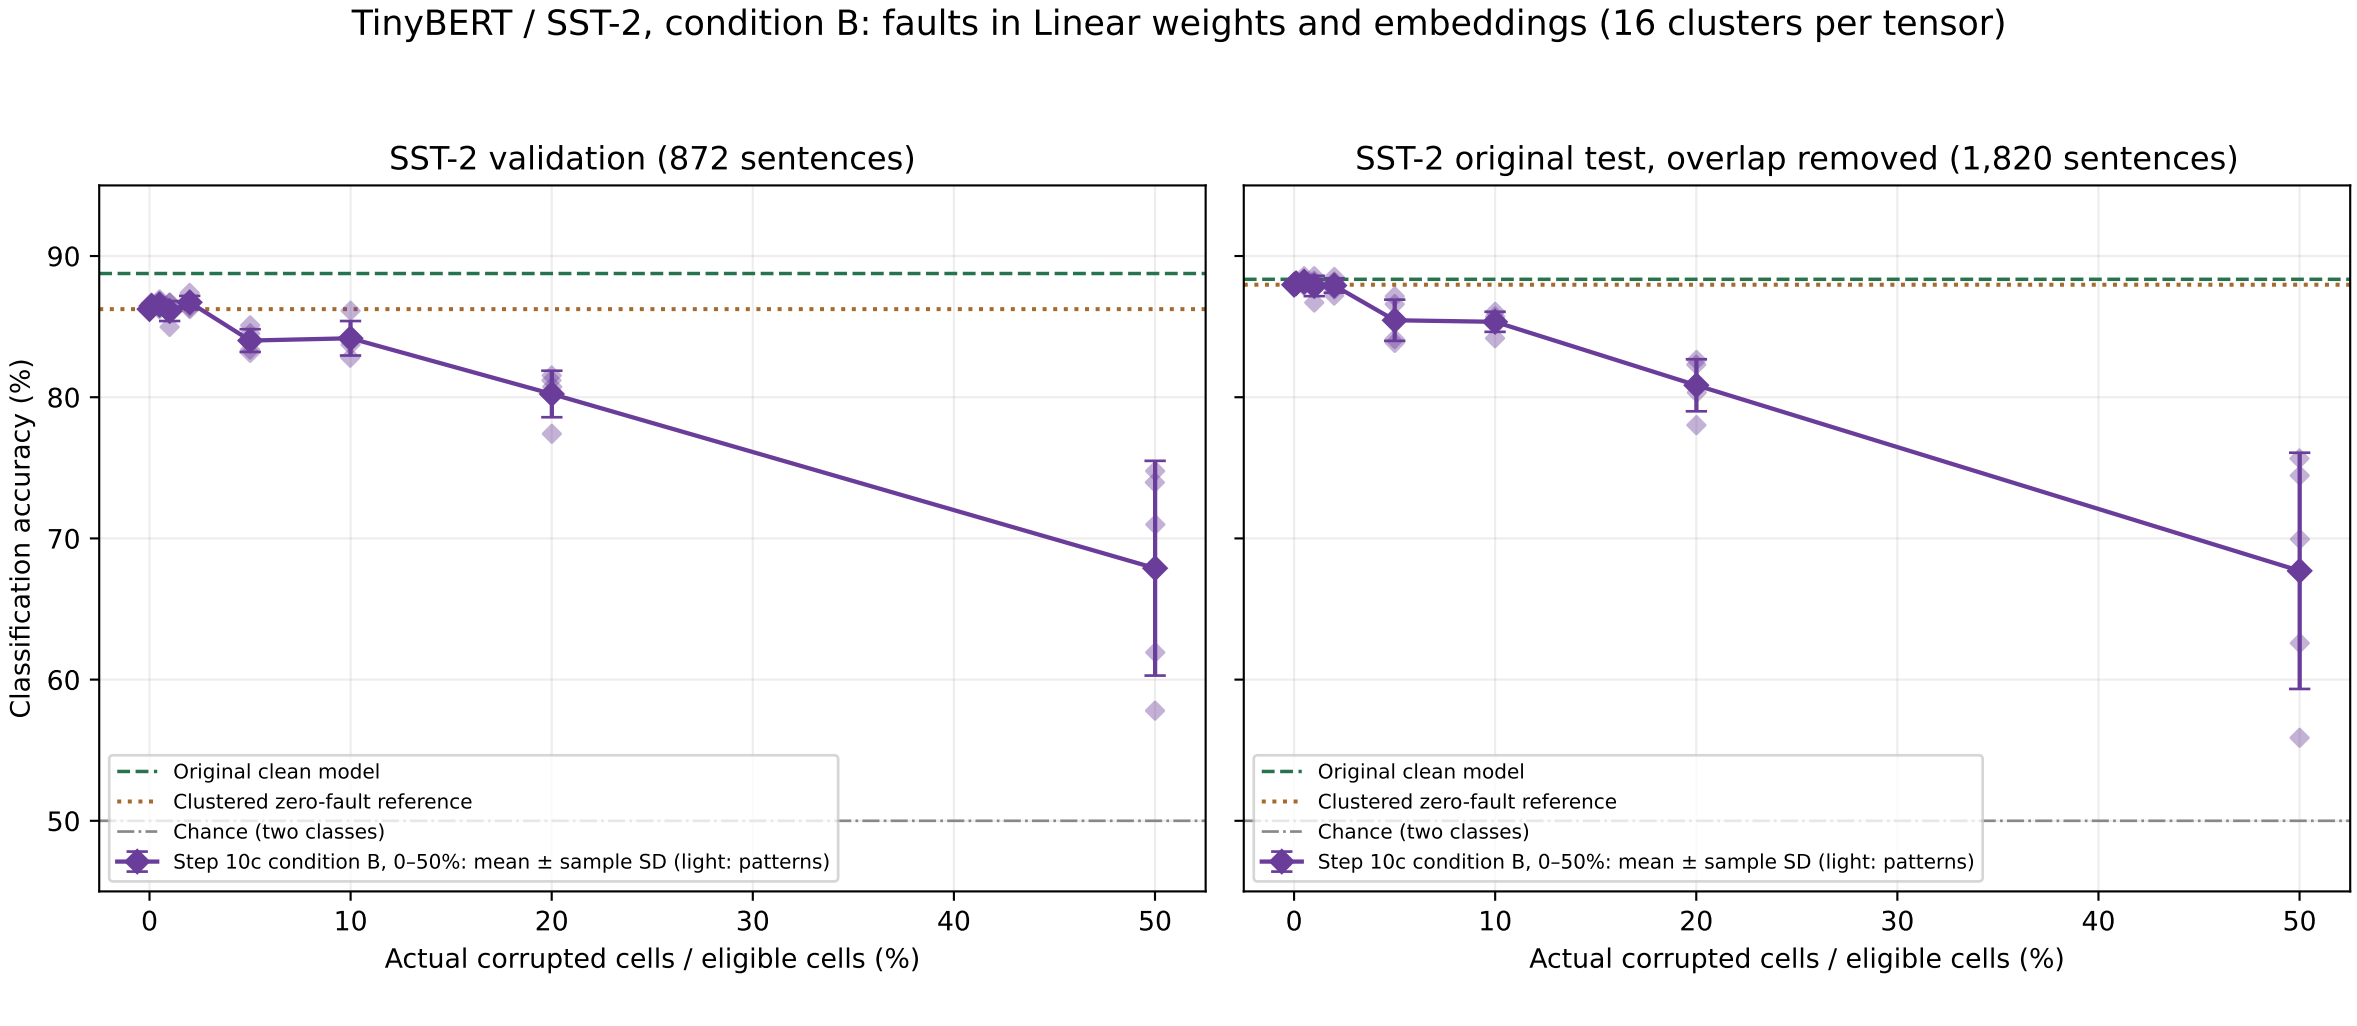
TinyBERT / SST-2, condition B: faults in Linear weights and embeddings (16 clusters per tensor)
SST-2 validation (872 sentences)
SST-2 original test, overlap removed (1,820 sentences)
Classification accuracy (%)
90
80
70
60
50
0
10
20
30
40
50
Actual corrupted cells / eligible cells (%)
Original clean model
Clustered zero-fault reference
Chance (two classes)
Step 10c condition B, 0–50%: mean ± sample SD (light: patterns)
[5] patterns per nonzero rate; zero faults: one reference. Embedding tables clustered and faulted; LayerNorm, biases, codebooks reliable.